

Reinforcement Learning

Notes

Jingye Wang

✉ wangjy5@shanghaitech.edu.cn

Spring 2020

Contents

1	Introduction	3
2	Review of Basic Probability	5
2.1	Interpretation of Probability	5
2.2	Transformations	5
2.3	Limit Theorem	5
2.4	Sampling & Monte Carlo Methods	6
2.5	Basic Inequalities	8
2.6	Concentration Inequalities	10
2.7	Conditional Expectation	12
3	Bandit Algorithms	14
3.1	Bandit Models	14
3.2	Stochastic Bandits	14
3.3	Greedy Algorithms	15
3.4	UCB Algorithms	16
3.5	Thompson Sampling Algorithms	17
3.6	Gradient Bandit Algorithms	18
4	Markov Chains	20
4.1	Markov Model	20
4.2	Basic Computations	20
4.3	Classifications	21

CONTENTS	2
4.4 Stationary Distribution	22
4.5 Reversibility	22
4.6 Markov chain Monte Carlo	23
5 Markov Decision Process	25
5.1 Markov Process	25
5.2 Markov Reward Process	25
5.3 Markov Decision Process	26
5.4 Dynamic Programming	28
6 Model-free Prediction	30
6.1 Monte-Carlo Policy Evaluation	30
6.2 Temporal-Difference Learning	31
7 Model-free Control	32
7.1 On Policy Monte-Carlo Control	33
7.2 On Policy Temporal-Difference Control	35
7.3 Off-Policy Q-Learning Control	36
7.4 Off-Policy Importance Sampling Control	36
8 Value Function Approximation	37
8.1 Introduction on Function Approximation	37
8.2 Incremental Method	38
8.3 Batch Methods	40
8.4 Deep Q-Learning	41
9 Policy Optimization	41
9.1 Policy Optimization	42
9.2 Monte-Carlo Policy Gradient	43
9.3 Actor-Critic Policy Gradient	46
9.4 Extension of Policy Gradient	47

6 Model-Free Prediction

Prediction: *Given an MDP and a policy π , find the corresponding value function.*

For an unknown MDP, we have no idea about its transition matrix and its state, action spaces, thus the methods we mentioned in previous sections fail to work. What we can do is interacting with the environment and collecting episodes.

Episodes: *An episode τ is a sequence of states and actions in the environment,*

$$\tau = (s_0, a_0, s_1, a_1, \dots).$$

We call an episode is complete if it ends with the terminal state.

In model free method, we only utilize the episode itself without further exploitation.

6.1 Monte-Carlo Policy Evaluation

Monte-Carlo (MC) policy evaluation methods learn directly from the *complete* episodes, thus it needs no knowledge of MDP transitions or rewards. The limitation of MC policy method methods is that it requires MDPs are episodic.

Monte-Carlo Policy Evaluation: *Given some complete episodes under the policy π , we can approximate the value of s by the average returns observed after visiting to s .*

Depending on when average returns for state s in an episode, there are two different implementations.

First-Visit Monte-Carlo Policy Evaluation: *Only at the first time-step t that state s is visited in an episodes, we do the following procedure:*

- Increment counter $N(s) \leftarrow N(s) + 1$;
- Increment the total return $return(s) \leftarrow return(s) + G_t$;
- Update the value by mean return $v(s) \leftarrow return(s)/N(s)$.

Every-Visit Monte-Carlo Policy Evaluation: *Every time-step t that state s is visited in an episode, we do the following procedure:*

- Increment counter $N(s) \leftarrow N(s) + 1$;
- Increment the total return $return(s) \leftarrow return(s) + G_t$;
- Update the value by mean return $v(s) \leftarrow return(s)/N(s)$.

Algorithm 10: First-Visit Monte-Carlo Policy Evaluation

```

1 Initialization:
2 the policy  $\pi$  to be evaluated;
3 initialize  $v(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$ ;
4 initialize  $return(s) \leftarrow$  an empty list for all  $s \in \mathcal{S}$ ;
5 Repeat:
6   Generate a complete episode  $\tau = (s_0, a_0, r_0, \dots, s_{T-1}, a_{T-1}, r_{T-1})$ ;
7    $G \leftarrow 0$ ;
8   For  $t = T - 1, T - 2, \dots, 0$ :
9      $G \leftarrow \gamma G + r_{t+1}$ ;
10    Unless  $s_t$  appears in  $s_0, s_1, \dots, s_{t-1}$ :
11      Append  $G$  to  $return(s_t)$ ;
12     $v(s_t) \leftarrow \text{average}(return(s_t))$ ;

```

Algorithm 11: Every-Visit Monte-Carlo Policy Evaluation

```

1 Initialization:
2 the policy  $\pi$  to be evaluated;
3 initialize  $v(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$ ;
4 initialize  $return(s) \leftarrow$  an empty list for all  $s \in \mathcal{S}$ ;
5 Repeat:
6   Generate a complete episode  $\tau = (s_0, a_0, r_0, \dots, s_{T-1}, a_{T-1}, r_{T-1})$ ;
7    $G \leftarrow 0$ ;
8   For  $t = T - 1, T - 2, \dots, 0$ :
9      $G \leftarrow \gamma G + r_{t+1}$ ;
10    Append  $G$  to  $return(s_t)$ ;
11     $v(s_t) \leftarrow \text{average}(return(s_t))$ ;

```

By law of large numbers, both of two methods can achieve $v(s) \rightarrow v^\pi(s)$ as $N(s) \rightarrow \infty$. In practice, we can use a trick *incremental mean* to simplify the calculation.

Differences between DP and MC for policy evaluation:

- DP computes v_k by bootstrapping the rest of the expected return calculated with v_{k-1} ;
- DP iterates on Bellman expectation backup:

$$v_k(s) \leftarrow \sum_{a \in \mathcal{A}} \pi(a|s) \left(R_s^a + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) v_{k-1}(s') \right).$$

- MC updates the empirical mean return with a sampled episode:

$$v(s_t) \leftarrow v(s_t) + \alpha(G_{k,t} - v(s_t)).$$

Advantages of MC over DP:

- MC can work when the environment is unknown;
- Working with sampled episodes has a huge advantage. Even with the complete knowledge of the environment's dynamics, the complexity still could be a challenge;
- Cost of estimating a single state's value is independent of the total number of states. So one can sample episodes starting from the states of interest then average returns.

6.2 Temporal-Difference Learning

Unlike MC methods, temporal-difference (TD) does not require the episodes are complete. TD methods can learn from incomplete episodes by bootstrapping.

Temporal-Difference Learning: *Given some incomplete episodes under the policy π , we update value $v(s_t)$ toward estimated return $r_{t+1} + \gamma v(s_{t+1})$:*

$$v(s_t) \leftarrow v(s_t) + \alpha(r_{t+1} + \gamma v(s_{t+1}) - v(s_t)),$$

where $r_{t+1} + \gamma v(s_{t+1})$ is called the TD target, α is the step-size, and we call

$$\delta_t = r_{t+1} + \gamma v(s_{t+1}) - v(s_t)$$

the TD error.

Algorithm 12: TD(0) Evaluation

```

1 Initialization:
2   the policy  $\pi$  to be evaluated;
3   the step size  $\alpha \in (0, 1]$ ;
4   initialize  $v(s) \in \mathbb{R}$  arbitrarily for all  $s \in \mathcal{S}$  except that  $v(\text{terminal}) = 0$ ;
5 Repeat:
6   Generate initial state  $s$ ;
7   While  $s$  is not terminal:
8      $a \leftarrow \pi(s)$ ;
9      $r, s' \leftarrow \text{environment}(s, a)$ ;
10     $v(s) \leftarrow v(s) + \alpha(r + \gamma v(s') - v(s))$ ;
11     $s \leftarrow s'$ ;
```

Differences between MC and TD for policy evaluation:

- TD can learn online after every step;
- MC must wait until end of episode before return is known;

- TD can learn from incomplete sequences;
- MC can only learn from complete sequences;
- TD works in continuing (non-terminating) environments;
- MC only works in episodic (terminating) environments;
- TD exploits Markov property, and it is efficient in Markov environments;
- MC does not exploit Markov property, thus it is relatively effective in non-Markov environments.

Bootstrapping: involves old values, it is something like in-place update.

Sampling: samples to get an expectation.

A summary of *bootstrapping* and *sampling* for DP, MC, and TD is shown in Table 2.

Algorithm	Bootstraps	Sampling
Dynamic Programming	✓	
Monte-Carlo		✓
Temporal-Difference	✓	✓

Table 2: *Bootstrapping* and *sampling* for DP, MC, and TD.

n-step TD methods: *Unlike the simplest TD method, in n-step TD methods, the updating rule of value $v(s_t)$ is*

$$v(s_t) \leftarrow v(s_t) + \alpha(G_t^{(n)} - v(s_t)),$$

where $G_t^{(n)}$ is the n -step return

$$G_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^n v(s_{t+n}).$$

Notice that with additional definition, we can generalize TD to MC when $n \rightarrow \infty$.

TD(λ) methods: *To make use of the information from all time-steps, we can use weight $(1 - \lambda)\lambda^{n-1}$ to average n -step returns over different n as*

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)},$$

thus the updating rule for the value becomes

$$v(s_t) \leftarrow v(s_t) + \alpha(G_t^\lambda - v(s_t)).$$